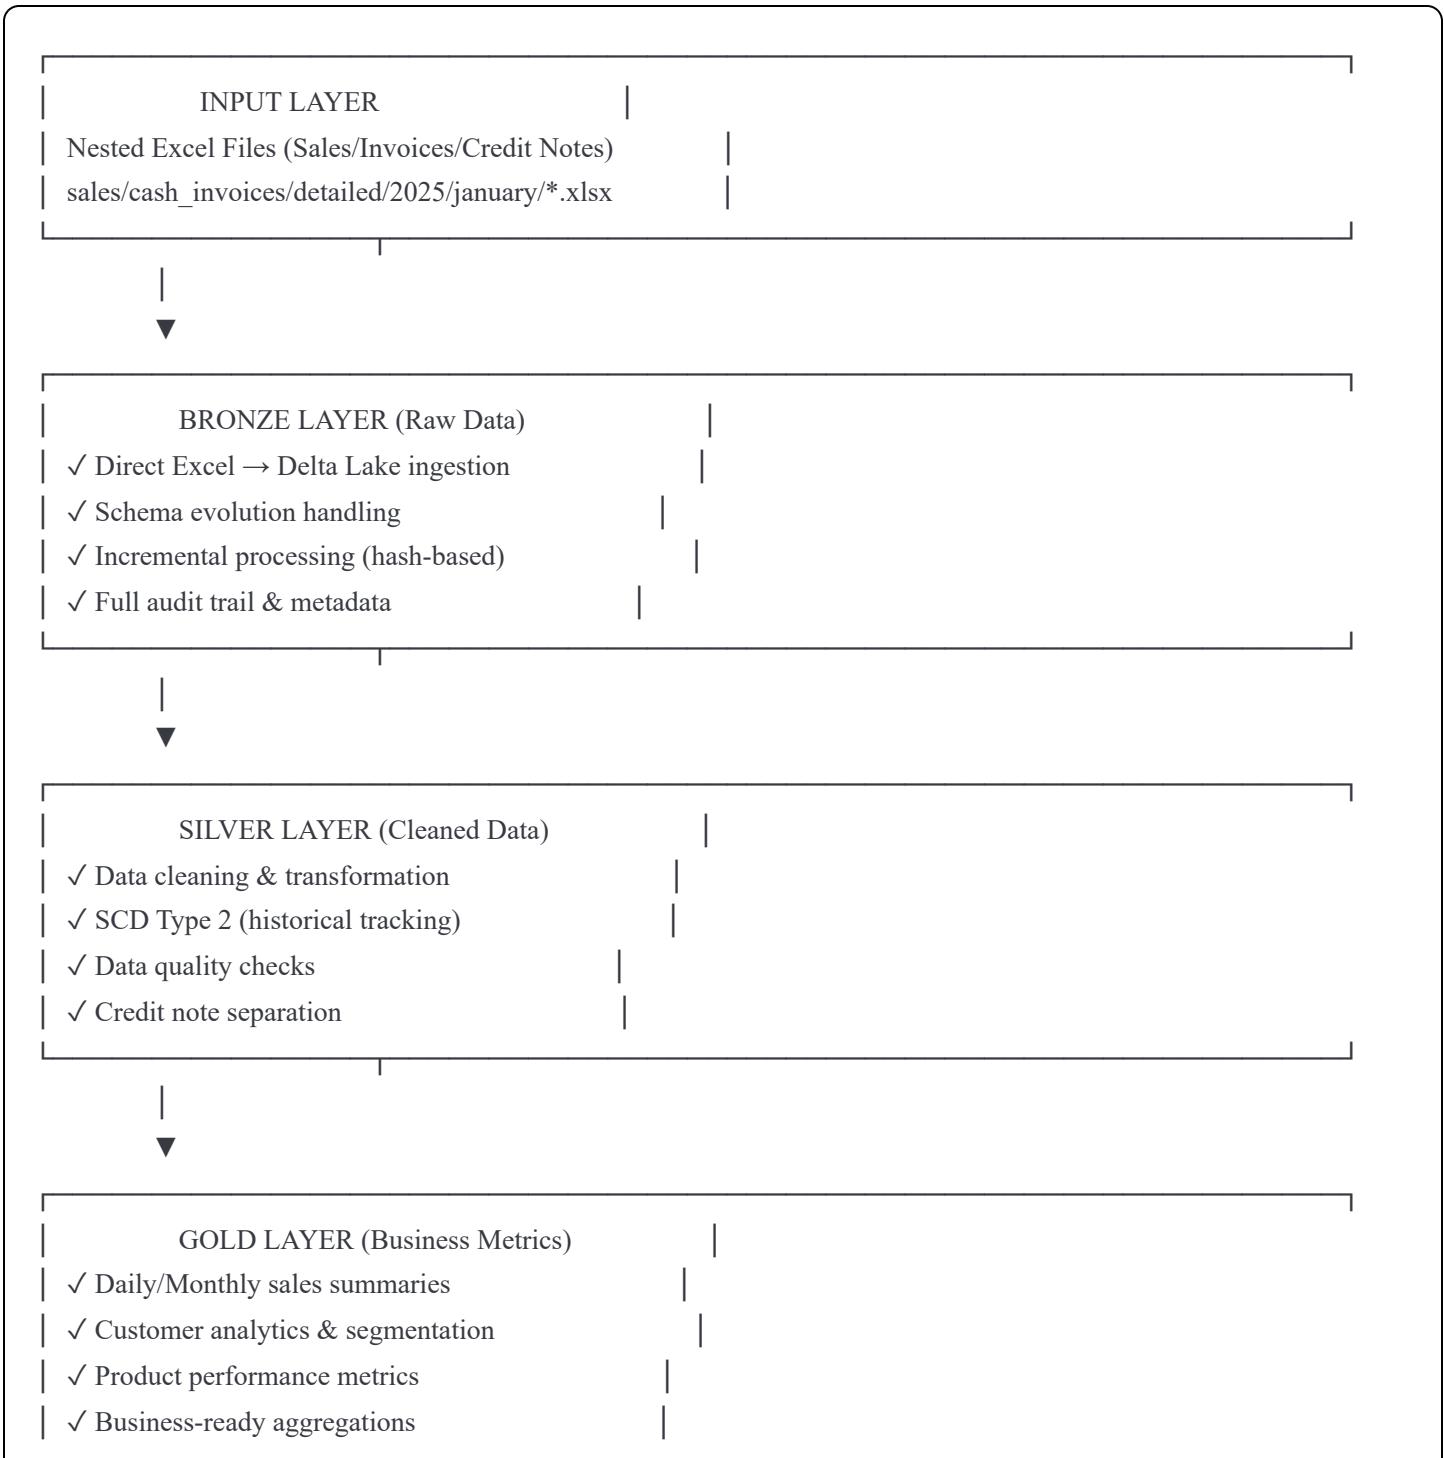


Medallion Architecture ETL Pipeline

Complete PySpark + Delta Lake + Prefect + PostgreSQL Solution

A production-ready, enterprise-grade ETL pipeline implementing the Medallion Architecture (Bronze-Silver-Gold) with comprehensive features including incremental processing, schema evolution, SCD Type 2, data quality checks, and full orchestration.

Architecture Overview



METADATA & MONITORING (PostgreSQL)

- ✓ Pipeline execution logs
- ✓ Data quality tracking
- ✓ Schema evolution history
- ✓ SCD2 operation logs

Key Features

Medallion Architecture

- **Bronze Layer:** Raw data ingestion from nested Excel files
- **Silver Layer:** Cleaned, transformed, and quality-checked data
- **Gold Layer:** Business-ready aggregations and metrics

Advanced Data Management

- **Incremental Processing:** Hash-based file change detection
- **Schema Evolution:** Automatic handling of schema changes
- **SCD Type 2:** Full historical tracking with validity periods
- **Data Partitioning:** Year/Month based partitioning for efficiency

Data Quality Framework

- Null value detection and flagging
- Duplicate detection (non-destructive)
- Amount validation checks
- Comprehensive quality logging

Production Features

- **Prefect Orchestration:** Workflow management and scheduling
- **PostgreSQL Metadata:** Complete audit trail and lineage
- **Error Handling:** Robust retry logic and error tracking

- **Monitoring:** Real-time pipeline and data quality monitoring
-

📁 Project Structure

```
medallion-etl-pipeline/
├── config.py          # Configuration management
├── metadata_logger.py # PostgreSQL logging
├── spark_utils.py     # Spark utilities & data quality
├── bronze_processor.py # Bronze layer processing
├── silver_processor.py # Silver layer with SCD2
├── main_pipeline.py   # Gold layer & Prefect orchestration
├── run_pipeline.py    # Main entry point
├── deploy_to_prefect.py # Prefect deployment
├── monitor_pipeline.py # Pipeline monitoring
├── requirements.txt    # Python dependencies
├── docker-compose.yml # PostgreSQL & Prefect setup
├── init_db.sql         # Database initialization
├── .env.example        # Environment variables template
└── README.md           # This file
```

🛠 Installation & Setup

Prerequisites

- Python 3.9+
- Docker & Docker Compose (for PostgreSQL & Prefect)
- Java 8 or 11 (for PySpark)

1. Clone Repository

```
bash
git clone <repository-url>
cd medallion-etl-pipeline
```

2. Setup Environment

Windows:

```
bash  
  
setup_environment.bat
```

Linux/Mac:

```
bash  
  
chmod +x setup_environment.sh  
./setup_environment.sh  
source venv/bin/activate
```

3. Configure Environment Variables

```
bash  
  
cp .env.example .env  
# Edit .env with your configuration
```

4. Start Infrastructure

```
bash  
  
# Start PostgreSQL and Prefect Server  
docker-compose up -d  
  
# Verify services are running  
docker-compose ps
```

5. Install Dependencies

```
bash  
  
pip install -r requirements.txt
```

🎯 Quick Start

Run Complete Pipeline

```
bash  
  
python run_pipeline.py
```

Run Specific Features

```
bash  
  
python run_pipeline.py --features sales_cash_invoices_detailed sales_credit_notes_detailed
```

Run Specific Layer

```
bash  
  
python run_pipeline.py --bronze-only  
python run_pipeline.py --silver-only  
python run_pipeline.py --gold-only
```

Pipeline Configuration

Configure Features (config.py)

```
python  
  
features: Dict[str, FeatureConfig] = {  
    "sales_cash_invoices_detailed": FeatureConfig(  
        name="sales_cash_invoices_detailed",  
        source_patterns=["sales/cash_invoices/detailed/**/*.*"],  
        partition_columns=["transaction_year", "transaction_month"],  
        scd2_enabled=True,  
        data_quality_checks=["null_check", "duplicate_check", "amount_check"]  
    )  
}
```

Configure Database (config.py or .env)

```
python  
  
POSTGRES_HOST=localhost  
POSTGRES_PORT=5432  
POSTGRES_DB=etl_metadata  
POSTGRES_USER=etl_user  
POSTGRES_PASSWORD=etl_password
```

Configure Spark (config.py)

```
python

spark_configs = {
    "spark.driver.memory": "4g",
    "spark.executor.memory": "4g",
    "spark.sql.shuffle.partitions": "200"
}
```

Prefect Orchestration

Deploy to Prefect

```
bash

python deploy_to_prefect.py
```

Access Prefect UI

```
http://localhost:4200
```

Run Deployed Pipeline

```
bash

prefect deployment run 'Medallion Architecture ETL Pipeline/medallion-etl-production'
```

Schedule Configuration

- Default: Daily at 2 AM UTC
- Modify in `deploy_to_prefect.py`

Monitoring & Metrics

View Pipeline Execution Summary

```
bash

python monitor_pipeline.py
```

Query Metadata Database

```
sql

-- Recent pipeline runs
SELECT * FROM etl_logs.pipeline_runs
ORDER BY started_at DESC LIMIT 10;

-- Data quality issues
SELECT * FROM etl_logs.data_quality_checks
WHERE check_status = 'FAILED';

-- Schema evolution history
SELECT * FROM etl_logs.schema_evolution
ORDER BY change_date DESC;

-- SCD2 operations
SELECT * FROM etl_logs.scd2_history
ORDER BY operation_timestamp DESC;
```

Access Prefect Monitoring

- **UI:** <http://localhost:4200>
 - **Flow Runs:** Monitor execution status
 - **Logs:** View detailed execution logs
 - **Metrics:** Track performance over time
-

📁 Data Organization

Input Structure

```
C:/etl/data/input/
├── sales/
│   ├── cash_invoices/
│   │   ├── detailed/
│   │   │   └── 2025/
│   │   │       └── january/
│   │   │           └── invoices_jan_2025.xlsx
│   │   └── summarized/
│   │       └── 2025/
```

```
| |    └── january/
|    ├── cash_sales/
|    |    ├── detailed/
|    |    └── summarized/
|    └── credit_notes/
|        ├── detailed/
|        └── summarized/
```

Output Structure (Delta Lake)

```
C:/etl/data/lakehouse/
├── bronze/
│   ├── bronze_sales_cash_invoices_detailed_sheet1/
│   ├── bronze_sales_cash_invoices_summarized_sheet1/
│   └── ...
├── silver/
│   ├── silver_sales_cash_invoices_detailed_invoices/
│   ├── silver_sales_cash_invoices_detailed_credits/
│   └── ...
└── gold/
    ├── daily_sales_summary/
    ├── monthly_sales_summary/
    ├── customer_analytics/
    └── product_analytics/
```

Data Quality Framework

Implemented Checks

1. **Null Checks:** Identify missing critical values
2. **Duplicate Detection:** Flag potential duplicates
3. **Amount Validation:** Verify numeric ranges
4. **Credit Note Validation:** Ensure proper sign

Quality Flags in Data

```
python
```

```
# Automatically added to dataframes
dq_null_flag_customer_code
dq_null_flag_amount
dq_duplicate_flag
```

Quality Reporting

All quality issues logged to PostgreSQL:

```
sql

SELECT
    table_name,
    check_type,
    check_status,
    records_affected,
    check_details
FROM etl_logs.data_quality_checks
WHERE check_status IN ('FAILED', 'WARNING');
```

SCD Type 2 Implementation

Features

- Automatic change detection via hash comparison
- Historical record preservation
- Current record flagging
- Version tracking

SCD2 Columns

```
python

_scd2_valid_from    # Timestamp when record became active
_scd2_valid_to      # Timestamp when record expired (NULL if current)
_scd2_is_current    # Boolean flag for current record
_scd2_version       # Version number
_business_key_hash  # Hash of business keys
_value_hash         # Hash of all values for change detection
```

Query Current Records

```
python  
  
df = spark.read.format("delta").load(table_path)  
current_df = df.filter(F.col("_scd2_is_current") == True)
```

Query Historical Records

```
python  
  
# Get all versions of a customer  
customer_history = df.filter(  
    (F.col("customer_code") == "CUST001") &  
    (F.col("_scd2_is_current") == False)  
).orderBy("_scd2_valid_from")
```

🔧 Advanced Configuration

Enable/Disable Features

```
python  
  
# In config.py  
enable_schema_evolution: bool = True  
enable_data_quality_checks: bool = True  
enable_scd2: bool = True
```

Tune Performance

```
python  
  
# Spark configurations  
spark_configs = {  
    "spark.sql.shuffle.partitions": "200", # Adjust based on data size  
    "spark.sql.adaptive.enabled": "true",  
    "spark.sql.adaptive.coalescePartitions.enabled": "true"  
}
```

Customize Partitioning

```
python
```

```
partition_columns: List[str] = ["transaction_year", "transaction_month"]
```

Troubleshooting

Common Issues

1. Spark Memory Errors

```
python  
  
# Increase memory in config.py  
driver_memory: str = "8g"  
executor_memory: str = "8g"
```

2. PostgreSQL Connection Failed

```
bash  
  
# Check if PostgreSQL is running  
docker-compose ps  
  
# Restart services  
docker-compose restart postgres
```

3. Delta Lake Write Errors

```
bash  
  
# Clear checkpoints  
rm -rf C:/etl/checkpoints/*  
  
# Vacuum old files  
spark.sql("VACUUM delta.`path/to/table` RETAIN 168 HOURS")
```

4. Excel Read Errors

```
bash  
  
# Install additional Excel libraries  
pip install xlrd openpyxl python-calamine
```

Best Practices

1. Incremental Processing

- Pipeline automatically detects new/modified files via hash comparison
- Only processes changed data
- Full audit trail maintained

2. Error Handling

- All errors logged to PostgreSQL
- Automatic retry for transient failures
- Partial success handling (some files can fail)

3. Data Quality

- Non-destructive checks (flags added, data preserved)
- Business rules monitored but not enforced
- Quality reports for manual review

4. Performance

- Partitioned Delta tables for fast queries
 - Adaptive query execution enabled
 - Parallel processing of features
-

Security Considerations

1. **Credentials:** Use environment variables, never hardcode
 2. **Database:** Use connection pooling and prepared statements
 3. **Files:** Implement file system permissions
 4. **Logs:** Sanitize sensitive data in logs
-

Production Deployment

Deployment Checklist

- Configure production database credentials
- Set appropriate Spark memory settings
- Configure Prefect schedules
- Set up monitoring alerts
- Enable backups for metadata database
- Configure Delta Lake retention policies
- Set up log rotation
- Test disaster recovery procedures

Scaling Considerations

- Use Spark cluster mode for large datasets
 - Implement Delta Lake optimize operations
 - Consider cloud storage (S3, Azure Blob, GCS)
 - Use distributed Prefect agents
-

Contributing

Contributions welcome! Please:

1. Fork the repository
 2. Create feature branch
 3. Add tests
 4. Submit pull request
-

License

MIT License - feel free to use in your projects

Support

For issues and questions:

- Create GitHub issue
 - Check documentation
 - Review metadata logs in PostgreSQL
-

Learning Resources

- [Delta Lake Documentation](#)
 - [PySpark Documentation](#)
 - [Prefect Documentation](#)
 - [Medallion Architecture](#)
-

Built with ❤️ for robust, scalable ETL pipelines